

## R

## Simulation and Optimization

R.M. Ripley

Department of Statistics  
University of Oxford

2008/9

## Uses of Simulation

- Statistical models are often mathematically intractable.
- Generate multiple samples from model by simulation.
- Use these samples to investigate the behaviour of the model.
- Need realizations of random variables with various different distributions, our ***random numbers***.
- Details of how to use the random numbers to create samples will not be considered here.

## Random number generation

- Random numbers calculated on a computer are not random.
- They are **pseudo-random**, following a predicted sequence, but in the short-term (i.e. anything but *very* long-term) will appear random.
- This is useful, as two sets of random numbers of the same size from the same generator using the same initial value will be exactly the same.
- Details of generation of random samples from standard distributions need not concern us, as R provides functions.

## Using Random Number Generators in R

- Each time a random number is required, R will use and update a variable called `.Random.seed` which is in your workspace.
- At the first use, if `.Random.seed` does not exist, one will be created, with a value generated from the time.
- The function `set.seed(n)` will set `.Random.seed` to a value derived from the argument `n`. Use `set.seed()` to repeat the same sequence of random numbers.
- It is possible to save and restore `.Random.seed` within your functions, but take care with *scope* (see (much) later!).

## R functions for generating random samples

To obtain a sample of size  $n$

<code>runif(n,min=0,max=1)</code>	Uniform
<code>rnorm(n,mean=0,sd=1)</code>	Normal
<code>rexp(n,rate=1)</code>	Exponential with mean 1/rate
<code>rt(n,df)</code>	t with df degrees of freedom
<code>rbinom(n,size,prob)</code>	Binomial, successes in <b>size</b> trials with probability of success <b>prob</b>

The parameters can usually be specified as vectors, so that non-iid samples can be obtained. For example,  
`contam<-rnorm(100,0,(1+2*rbinom(100,1,0.05)))`  
 will generate 100 samples from the contaminated normal distribution in which a sample is from  $N(0, 1)$  with probability 0.95 and otherwise from  $N(0, 9)$ .

## Further sampling

The function `sample` re-samples from a data vector, with or without replacement. It has several forms:

<code>sample(n)</code>	select a random perm. from $1, \dots, n$
<code>sample(x)</code>	randomly permute <b>x</b> , for <code>length(x) &gt; 1</code>
<code>sample(x, repl=T)</code>	a bootstrap sample
<code>sample(x, n)</code>	sample <b>n</b> items from <b>x</b> w/o replacement
<code>sample(x, n, repl=T)</code>	sample <b>n</b> items from <b>x</b> with replacement
<code>sample(x, n, repl=T, prob=probs)</code>	probability sample of <b>n</b> items from <b>x</b> .

## The Optimization Problem

- Given a function  $f(\mathbf{x})$ , what value of  $\mathbf{x}$  makes  $f(\mathbf{x})$  as small or as large as possible?
- In a statistical context,  $\mathbf{x}$  will usually be the **parameters** of a model, and  $f(\mathbf{x})$  either the model likelihood to be maximized or some measure of discrepancy between data and predictions to be minimized
- The optimal set of parameters will give the *best fit*.
- Only need to consider *small* as  $-f(\mathbf{x})$  is *large* when  $f(\mathbf{x})$  is *small*.
- We consider here **general-purpose optimizers**.

## Local and Global Minima

- The (negative of the) likelihood for the General Linear Model (and that for many other linear models) is well-behaved: it has a single, global minimum.
- For more complicated models there may be many local minima.
- Finding a global minimum is difficult, and not always important. Only if local minima are widely separated in parameter space are they likely to invalidate our conclusions.
- We will concentrate on methods of finding local minima.
- Check for different local minima by altering the initial values, algorithm used, or other parameters of the fitting process.

## Univariate Optimization

**optimize** (or **optimise**) finds a (possibly local) minimum of a function in a specified interval with respect to its first argument.

- Function to be minimized is first argument of **optimize**.
- Can pre-specify the function or include it in the command (cf. panel functions)

### Example

```
f <- function (x,a) (x-a)^2
xmin<- optimize(f, interval=c(0, 1), a = 1/3)
or
xmin<- optimize(function(x,a) (x-a)^2,
                 interval= c(0, 1), a = 1/3)
```

Note how the (fixed) parameter **a** is passed into **f**.

## Other optimize() Arguments

- Interval within which to search can be specified by **interval=** or **upper=**, **lower=**.
- To *maximize*, set **maximum=TRUE**.
- Accuracy can be set using the **tol=** argument.
- Note the order of arguments:  
**optimize(f, interval, ..., lower, upper, maximum, tol)**
- The ... can be named or unnamed and will be passed to **f**.
- Arguments after the ... must be specified by name.
- **optimize** returns a list with two items:
  - minimum The value of **x** at which **f(x)** is minimized.
  - objective The value of **f(x)** at **x=minimum**

## General Optimization

- In more than one dimension the problem is harder.
- R has several different functions: most flexible is **optim()** which includes several different algorithms.
- Algorithm of choice depends on how easy it is to calculate derivatives for the function. Usually better to supply a function to calculate derivatives, but may be unnecessary extra work.
- Ensure the problem is scaled so that unit change in any parameter gives approximately unit change in objective.
- Experiment with different methods...

## optim(): Nelder-Mead method

- The **default** method
- Basic idea: for a function with **n** vertices, choose a polygon with **n+1** vertices. At each step, alter vertex with minimum **f(x)** to improve the objective function, by *reflection*, *expansion* or *contraction*
- Does not use derivative information
- Useful for non-differentiable functions
- May be rather slow

## optim(): BFGS method

- A **quasi-Newton** method: builds up approximation to Hessian matrix from gradients at start and finish of steps
- Uses the approximation to choose new search direction
- Performs line search in this direction
- Update term for the Hessian approximation is due to Broyden, Fletcher, Goldfarb and Shanno (proposed separately by all four in 1970)
- Uses derivative information, calculated either from a user-supplied function or by finite differences
- If dimension is large, the matrix stored may be very large

## optim(): CG method

- A **conjugate gradient** method: chooses successive search directions that are analogous to axes of an ellipse
- Does not store a Hessian matrix
- Three different formulae for the search directions are implemented: *Fletcher-Reeves*, *Polak-Ribiere* or *Beale-Sorenson*
- Less robust than BFGS method
- Uses derivative information, calculated either from a user-supplied function or by finite differences

## optim(): L-BFGS-B method

- A **limited memory version of BFGS**
- Does not store a Hessian matrix, only a limited number of update steps for it
- Uses derivative information, calculated either from a user-supplied function or by finite differences
- Can restrict the solution to lie within a "box", the only method of `optim()` that can do this

## optim(): SANN method

- A variant of **simulated annealing**
- A stochastic algorithm
- Accepts changes which increase the objective with positive probability (when minimising!)
- Does not use derivative information
- Can be very slow to converge but may find a 'good' solution quickly

## How to use optim()

```
optim(par, fn, gr=NULL, ...,  
method=c("Nelder-Mead", "BFGS", "CG", "L-BFGS-B", "SANN"),  
lower=-Inf, upper=Inf, control=list(), hessian=FALSE)
```

**par** Starting values of the parameters.  
**fn** The function (supply as for **optimize**)  
**gr** Function to calculate the derivative. Only relevant for methods "BFGS", "CG", or "L-BFGS-B".  
 ... Other parameters for (both) **fn** and **gr**.  
**method** Algorithm to use  
**lower, upper** Vectors of limits for parameters if required. Only allowed if **method**='L-BFGS-B'.  
**control** control options: see next slide  
**hessian** Logical: whether to return a hessian estimate calculated by finite differences.

## optim(): Control Options

There are very many. The most important are:

- trace** A positive integer: higher values give more information
- fnscale** An overall scaling: if negative, maximization will be performed
- parscale** A vector of scalings for the parameters
- maxit** Maximum number of iterations to be performed: may be useful to terminate unsuccessful attempts  
Also used to perform one or two steps if convergence is unimportant
- type** Used to select formula for "CG" method

## optim(): Components of Return Value

**par** best set of parameters  
**value** value of **fn** corresponding to **par**  
**counts** number of calls to **fn** and **gr**: excludes calls for purposes of approximating derivatives or Hessian  
**convergence** error code: 0=success, 1=max iterations reached, 10= degeneracy of Nelder-Mead simplex, 51=warning from "L-BFGS-B", 52=error from "L-BFGS-B"  
**message** further information, if any  
**hessian** If requested, a symmetric matrix estimate of the Hessian at the solution: Hessian of unconstrained problem, even if constraints are active

## Constrained Optimization

- Often possible to constrain parameters to be positive by transforming to logarithmic scale
- **optim()** with **method**="L-BFGS-B" will accept box constraints
- **nlmminb** is an alternative function allowing box constraints
- **constrOptim** is a wrapper for **optim** which enforces linear inequality constraints by adding a barrier function: can use with any **optim** method except "L-BFGS-B"

## Arguments of constrOptim

theta starting values of parameters  
 f function  
 grad gradient. Must be given if using a method which needs derivatives.  
 ui, ci constraints are that `ui %% theta - ci >= 0`  
 mu tuning parameter which takes small values  
 control as for `optim`  
 method as for `optim`  
 outer.iterations Number of iterations in outer loop  
 outer.eps Accuracy required in outer loop  
 ... other parameters for your function  
 Returns values as for `optim` plus value of barrier function and number of outer iterations

## Exercises 5 Page 1

- In the package **survival** there is a dataset called **ovarian**. Your task is to fit a log normal model to relate survival time to age, and treatment, using **optim**.
- The likelihood to be used is

$$\prod_{\delta_i=1} f(t_i) \prod_{\delta_i=0} (1 - F(t_i))$$

where  $\delta_i$  is the censoring indicator, `ovarian$fustat`,  $f(t_i)$  is the density and  $F(t_i)$  the cumulative distribution. The mean to be fitted, on a log scale, will be  $\mathbf{x}\beta$  where  $\mathbf{x}$  is the covariate matrix with a column of 1's added for the mean, and  $\beta$  are parameters to be estimated. A common variance should be estimated.

## Exercises 5 Page 2

- Write a function to calculate the likelihood from the parameters, and maximize it using **optim**. Then write a function to calculate the derivatives and use that as well. Try different methods of fitting.
- Hints:
  - Transform  $\sigma$  to a log scale to avoid constraints.
  - Start with just the mean and no covariates.
  - For `method='BFGS'`, you may need to use **parscale**. If you know the approximate gradients, set **parscale** proportional to these. Even if not essential it may speed things up.
  - After writing the function to calculate the derivatives, check it using finite differences and the function which calculates the likelihood.
- Check your results using **survreg**.